

```
In [7]: # Import necessary Libraries
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
import time

# Define data transformations
# Converting images to tensors and normalizing them to range [-1, 1]
transform = transforms.Compose([
    transforms.ToTensor(), # Convert PIL images to tensors (0-1 range)
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)) # Normalize RGB channel
])

# Load and prepare the CIFAR-10 dataset
# Training set with 50,000 images
train_dataset = datasets.CIFAR10(root='./data', train=True, download=True, transform=transform)
# Test set with 10,000 images
test_dataset = datasets.CIFAR10(root='./data', train=False, download=True, transform=transform)

# Create data Loaders for batch processing
# Batch size of 64 is a good balance between memory usage and training speed
train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=64, shuffle=False)

# Problem 1a: Multi-Layer Perceptron Implementation
class MLP(nn.Module):
    def __init__(self):
        super(MLP, self).__init__()
        # Sequential container for the network layers
        self.model = nn.Sequential(
            nn.Flatten(), # Flatten 32x32x3 input images to 3072 vector
            # First hidden Layer: 3072 -> 256 neurons
            nn.Linear(32 * 32 * 3, 256),
            nn.ReLU(), # ReLU activation for non-linearity
            # Second hidden Layer: 256 -> 128 neurons
            nn.Linear(256, 128),
            nn.ReLU(),
            # Third hidden Layer: 128 -> 64 neurons
            nn.Linear(128, 64),
            nn.ReLU(),
            # Output Layer: 64 -> 10 classes (CIFAR-10 classes)
            nn.Linear(64, 10)
        )

        def forward(self, x):
            return self.model(x)

# Set up device (GPU if available, else CPU)
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

```
model = MLP().to(device)

# Define loss function and optimizer
criterion = nn.CrossEntropyLoss() # Appropriate for multi-class classification
optimizer = optim.Adam(model.parameters(), lr=0.001) # Adam optimizer with Learnin

# Training hyperparameters
epochs = 20
# Lists to store metrics for plotting
train_loss_history, val_loss_history = [], []
train_acc_history, val_acc_history = [], []

print("Training started...")
start_time = time.time()

# Training Loop
for epoch in range(epochs):
    # Training phase
    model.train()
    running_loss = 0.0
    correct_train = 0
    total_train = 0

    for images, labels in train_loader:
        # Move data to appropriate device
        images, labels = images.to(device), labels.to(device)

        # Forward pass
        outputs = model(images)
        loss = criterion(outputs, labels)

        # Backward pass and optimization
        optimizer.zero_grad() # Clear previous gradients
        loss.backward() # Compute gradients
        optimizer.step() # Update weights

        # Calculate training metrics
        running_loss += loss.item()
        _, predicted = outputs.max(1)
        correct_train += (predicted == labels).sum().item()
        total_train += labels.size(0)

    # Calculate epoch metrics for training
    train_loss = running_loss / len(train_loader)
    train_acc = 100 * correct_train / total_train
    train_loss_history.append(train_loss)
    train_acc_history.append(train_acc)

    # Validation phase
    model.eval() # Set model to evaluation mode
    val_loss = 0.0
    correct_val = 0
    total_val = 0

    # No gradient computation needed for validation
    with torch.no_grad():


```

```
        for images, labels in test_loader:
            images, labels = images.to(device), labels.to(device)
            outputs = model(images)
            loss = criterion(outputs, labels)
            val_loss += loss.item()
            _, predicted = outputs.max(1)
            correct_val += (predicted == labels).sum().item()
            total_val += labels.size(0)

        # Calculate epoch metrics for validation
        val_loss /= len(test_loader)
        val_acc = 100 * correct_val / total_val
        val_loss_history.append(val_loss)
        val_acc_history.append(val_acc)

    # Print epoch results
    print(f"Epoch [{epoch + 1}/{epochs}], Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.2f}, Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.2f}")

    # Print total training time
    end_time = time.time()
    print(f"Training completed in {end_time - start_time:.2f} seconds.")

# Save the trained model
torch.save(model.state_dict(), "mlp_cifar10.pth")

# Plotting training and validation metrics
plt.figure(figsize=(12, 5))

# Loss plot
plt.subplot(1, 2, 1)
plt.plot(train_loss_history, label='Training Loss')
plt.plot(val_loss_history, label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.title('Loss Over Epochs')

# Accuracy plot
plt.subplot(1, 2, 2)
plt.plot(train_acc_history, label='Training Accuracy')
plt.plot(val_acc_history, label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.title('Accuracy Over Epochs')

plt.show()

# Final model evaluation
model.eval()
y_true, y_pred = [], []

# Collect predictions on test set
with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
```

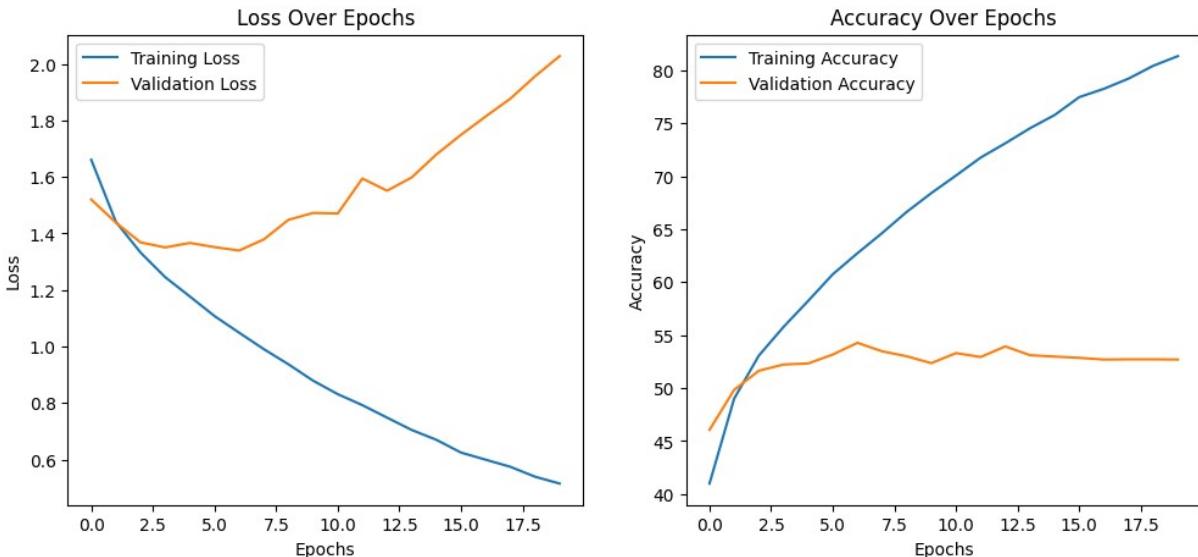
```
_ , predicted = outputs.max(1)
y_true.extend(labels.cpu().numpy())
y_pred.extend(predicted.cpu().numpy())

# Generate and print classification metrics
print("\nClassification Report:")
print(classification_report(y_true, y_pred, target_names=train_dataset.classes))

# Generate and plot confusion matrix
print("\nConfusion Matrix:")
conf_matrix = confusion_matrix(y_true, y_pred)

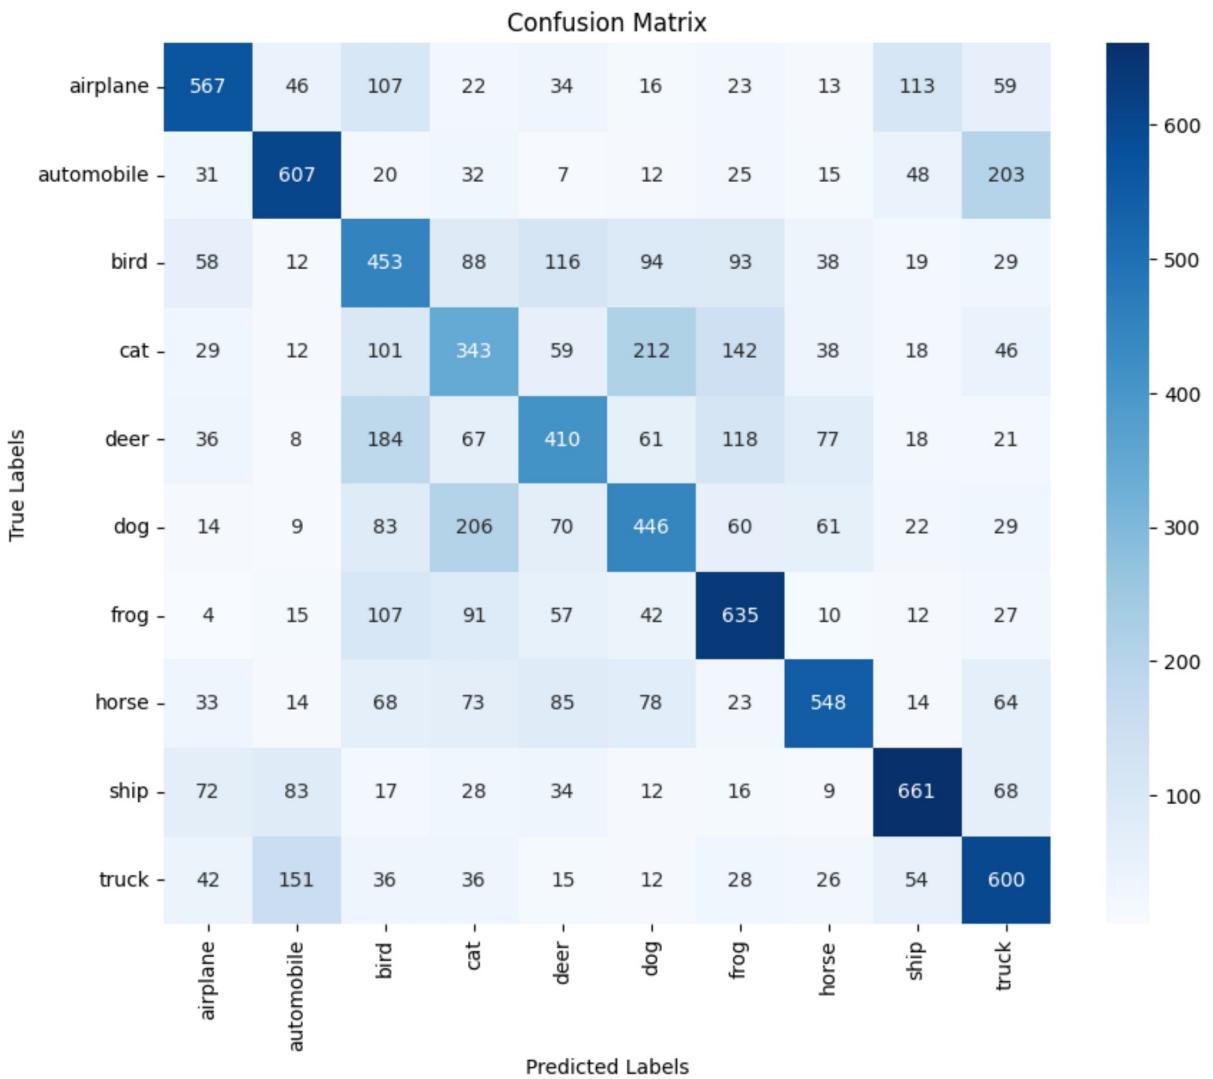
plt.figure(figsize=(10, 8))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=train_dataset.classes,
            yticklabels=train_dataset.classes)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix')
plt.show()
```

Files already downloaded and verified  
Files already downloaded and verified  
Training started...  
Epoch [1/20], Train Loss: 1.6607, Train Acc: 40.99%, Val Loss: 1.5201, Val Acc: 46.06%  
Epoch [2/20], Train Loss: 1.4410, Train Acc: 49.01%, Val Loss: 1.4383, Val Acc: 49.84%  
Epoch [3/20], Train Loss: 1.3324, Train Acc: 53.07%, Val Loss: 1.3684, Val Acc: 51.64%  
Epoch [4/20], Train Loss: 1.2453, Train Acc: 55.77%, Val Loss: 1.3506, Val Acc: 52.23%  
Epoch [5/20], Train Loss: 1.1770, Train Acc: 58.23%, Val Loss: 1.3664, Val Acc: 52.33%  
Epoch [6/20], Train Loss: 1.1077, Train Acc: 60.75%, Val Loss: 1.3517, Val Acc: 53.17%  
Epoch [7/20], Train Loss: 1.0488, Train Acc: 62.74%, Val Loss: 1.3396, Val Acc: 54.28%  
Epoch [8/20], Train Loss: 0.9907, Train Acc: 64.64%, Val Loss: 1.3789, Val Acc: 53.48%  
Epoch [9/20], Train Loss: 0.9369, Train Acc: 66.64%, Val Loss: 1.4483, Val Acc: 53.01%  
Epoch [10/20], Train Loss: 0.8790, Train Acc: 68.42%, Val Loss: 1.4725, Val Acc: 52.36%  
Epoch [11/20], Train Loss: 0.8313, Train Acc: 70.07%, Val Loss: 1.4707, Val Acc: 53.31%  
Epoch [12/20], Train Loss: 0.7923, Train Acc: 71.77%, Val Loss: 1.5945, Val Acc: 52.94%  
Epoch [13/20], Train Loss: 0.7482, Train Acc: 73.12%, Val Loss: 1.5512, Val Acc: 53.94%  
Epoch [14/20], Train Loss: 0.7046, Train Acc: 74.54%, Val Loss: 1.5983, Val Acc: 53.11%  
Epoch [15/20], Train Loss: 0.6698, Train Acc: 75.79%, Val Loss: 1.6795, Val Acc: 52.98%  
Epoch [16/20], Train Loss: 0.6245, Train Acc: 77.48%, Val Loss: 1.7491, Val Acc: 52.86%  
Epoch [17/20], Train Loss: 0.5995, Train Acc: 78.26%, Val Loss: 1.8140, Val Acc: 52.70%  
Epoch [18/20], Train Loss: 0.5743, Train Acc: 79.23%, Val Loss: 1.8769, Val Acc: 52.72%  
Epoch [19/20], Train Loss: 0.5393, Train Acc: 80.43%, Val Loss: 1.9558, Val Acc: 52.72%  
Epoch [20/20], Train Loss: 0.5150, Train Acc: 81.35%, Val Loss: 2.0275, Val Acc: 52.70%  
Training completed in 145.18 seconds.

**Classification Report:**

	precision	recall	f1-score	support
airplane	0.64	0.57	0.60	1000
automobile	0.63	0.61	0.62	1000
bird	0.39	0.45	0.42	1000
cat	0.35	0.34	0.35	1000
deer	0.46	0.41	0.43	1000
dog	0.45	0.45	0.45	1000
frog	0.55	0.64	0.59	1000
horse	0.66	0.55	0.60	1000
ship	0.68	0.66	0.67	1000
truck	0.52	0.60	0.56	1000
accuracy			0.53	10000
macro avg	0.53	0.53	0.53	10000
weighted avg	0.53	0.53	0.53	10000

**Confusion Matrix:**



In [2]:

```
#Problem 1b
class ComplexMLP(nn.Module):
    def __init__(self):
        super(ComplexMLP, self).__init__()
        self.model = nn.Sequential(
            nn.Flatten(),
            nn.Linear(32 * 32 * 3, 512),
            nn.ReLU(),
            nn.Linear(512, 256),
            nn.ReLU(),
            nn.Linear(256, 128),
            nn.ReLU(),
            nn.Linear(128, 64),
            nn.ReLU(),
            nn.Linear(64, 10)
        )

    def forward(self, x):
        return self.model(x)

complex_model = ComplexMLP().to(device)
optimizer = optim.Adam(complex_model.parameters(), lr=0.001)
criterion = nn.CrossEntropyLoss()
```

```
epochs = 20
complex_train_loss_history, complex_val_loss_history = [], []
complex_train_acc_history, complex_val_acc_history = [], []

print("Training started...")
start_time = time.time()

for epoch in range(epochs):
    complex_model.train()
    running_loss = 0.0
    correct_train = 0
    total_train = 0

    for images, labels in train_loader:
        images, labels = images.to(device), labels.to(device)

        #Forward pass
        outputs = complex_model(images)
        loss = criterion(outputs, labels)

        #Backward pass
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        running_loss += loss.item()
        _, predicted = outputs.max(1)
        correct_train += (predicted == labels).sum().item()
        total_train += labels.size(0)

    train_loss = running_loss / len(train_loader)
    train_acc = 100 * correct_train / total_train
    complex_train_loss_history.append(train_loss)
    complex_train_acc_history.append(train_acc)

    complex_model.eval()
    val_loss = 0.0
    correct_val = 0
    total_val = 0

    with torch.no_grad():
        for images, labels in test_loader:
            images, labels = images.to(device), labels.to(device)
            outputs = complex_model(images)
            loss = criterion(outputs, labels)
            val_loss += loss.item()
            _, predicted = outputs.max(1)
            correct_val += (predicted == labels).sum().item()
            total_val += labels.size(0)

    val_loss /= len(test_loader)
    val_acc = 100 * correct_val / total_val
    complex_val_loss_history.append(val_loss)
    complex_val_acc_history.append(val_acc)
```

```
print(f"Epoch [{epoch + 1}/{epochs}], Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.4f}\n")
end_time = time.time()
print(f"Training for Problem 1b completed in {end_time - start_time:.2f} seconds.\n")

plt.figure(figsize=(12, 5))
#Loss
plt.subplot(1, 2, 1)
plt.plot(complex_train_loss_history, label='Training Loss')
plt.plot(complex_val_loss_history, label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.title('Loss Over Epochs (Complex Model)')

#Accuracy
plt.subplot(1, 2, 2)
plt.plot(complex_train_acc_history, label='Training Accuracy')
plt.plot(complex_val_acc_history, label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.title('Accuracy Over Epochs (Complex Model)')

plt.show()

complex_model.eval()
y_true, y_pred = [], []

with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        outputs = complex_model(images)
        _, predicted = outputs.max(1)
        y_true.extend(labels.cpu().numpy())
        y_pred.extend(predicted.cpu().numpy())

#confusion matrix
print("\nClassification Report (Problem 1b):")
print(classification_report(y_true, y_pred, target_names=train_dataset.classes))

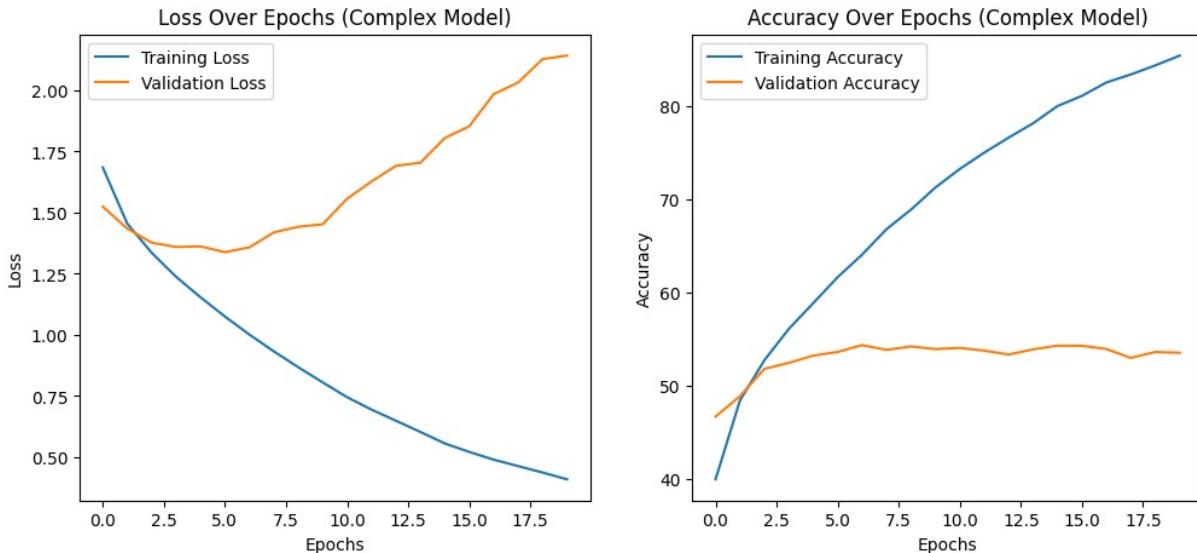
print("\nConfusion Matrix (Problem 1b):")
conf_matrix = confusion_matrix(y_true, y_pred)

plt.figure(figsize=(10, 8))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=train_data['label'].unique(), yticklabels=train_data['label'].unique())
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix (Problem 1b)')
plt.show()
```

Training started...

Epoch [1/20], Train Loss: 1.6847, Train Acc: 39.91%, Val Loss: 1.5249, Val Acc: 46.65%  
Epoch [2/20], Train Loss: 1.4554, Train Acc: 48.39%, Val Loss: 1.4352, Val Acc: 48.85%  
Epoch [3/20], Train Loss: 1.3359, Train Acc: 52.71%, Val Loss: 1.3772, Val Acc: 51.78%  
Epoch [4/20], Train Loss: 1.2378, Train Acc: 56.08%, Val Loss: 1.3596, Val Acc: 52.43%  
Epoch [5/20], Train Loss: 1.1539, Train Acc: 58.85%, Val Loss: 1.3621, Val Acc: 53.21%  
Epoch [6/20], Train Loss: 1.0752, Train Acc: 61.64%, Val Loss: 1.3383, Val Acc: 53.60%  
Epoch [7/20], Train Loss: 1.0014, Train Acc: 64.07%, Val Loss: 1.3580, Val Acc: 54.34%  
Epoch [8/20], Train Loss: 0.9328, Train Acc: 66.81%, Val Loss: 1.4192, Val Acc: 53.82%  
Epoch [9/20], Train Loss: 0.8679, Train Acc: 68.90%, Val Loss: 1.4424, Val Acc: 54.19%  
Epoch [10/20], Train Loss: 0.8060, Train Acc: 71.28%, Val Loss: 1.4523, Val Acc: 53.91%  
Epoch [11/20], Train Loss: 0.7453, Train Acc: 73.28%, Val Loss: 1.5571, Val Acc: 54.03%  
Epoch [12/20], Train Loss: 0.6944, Train Acc: 75.03%, Val Loss: 1.6275, Val Acc: 53.74%  
Epoch [13/20], Train Loss: 0.6489, Train Acc: 76.65%, Val Loss: 1.6915, Val Acc: 53.32%  
Epoch [14/20], Train Loss: 0.6031, Train Acc: 78.19%, Val Loss: 1.7041, Val Acc: 53.88%  
Epoch [15/20], Train Loss: 0.5561, Train Acc: 80.04%, Val Loss: 1.8051, Val Acc: 54.28%  
Epoch [16/20], Train Loss: 0.5215, Train Acc: 81.13%, Val Loss: 1.8531, Val Acc: 54.26%  
Epoch [17/20], Train Loss: 0.4898, Train Acc: 82.58%, Val Loss: 1.9845, Val Acc: 53.92%  
Epoch [18/20], Train Loss: 0.4634, Train Acc: 83.43%, Val Loss: 2.0323, Val Acc: 52.96%  
Epoch [19/20], Train Loss: 0.4373, Train Acc: 84.40%, Val Loss: 2.1279, Val Acc: 53.60%  
Epoch [20/20], Train Loss: 0.4094, Train Acc: 85.46%, Val Loss: 2.1423, Val Acc: 53.51%

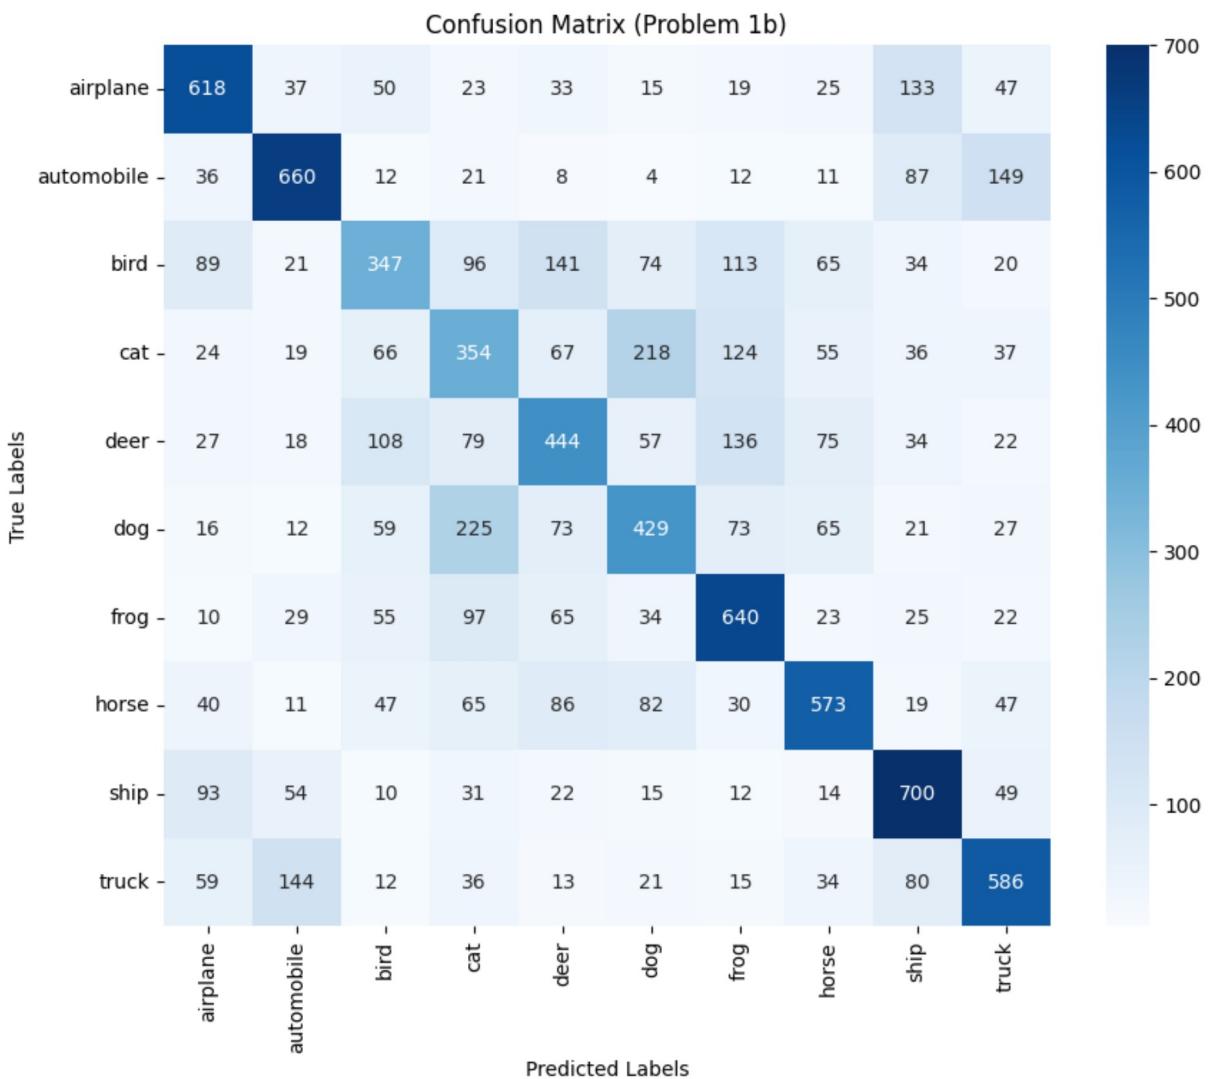
Training for Problem 1b completed in 149.73 seconds.



Classification Report (Problem 1b):

	precision	recall	f1-score	support
airplane	0.61	0.62	0.61	1000
automobile	0.66	0.66	0.66	1000
bird	0.45	0.35	0.39	1000
cat	0.34	0.35	0.35	1000
deer	0.47	0.44	0.45	1000
dog	0.45	0.43	0.44	1000
frog	0.55	0.64	0.59	1000
horse	0.61	0.57	0.59	1000
ship	0.60	0.70	0.65	1000
truck	0.58	0.59	0.58	1000
accuracy			0.54	10000
macro avg	0.53	0.54	0.53	10000
weighted avg	0.53	0.54	0.53	10000

Confusion Matrix (Problem 1b):



In [3]:

```
#Problem 2a
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import time

#dataset
url = "https://raw.githubusercontent.com/HamedTabkhi/Intro-to-ML/refs/heads/main/Da
data = pd.read_csv(url)

for col in data.select_dtypes(include=['object']).columns:
    data[col] = data[col].astype('category').cat.codes

target_column = data.columns[-1]
features = data.drop([target_column], axis=1)
targets = data[target_column].values.reshape(-1, 1)
```

```
features = features.apply(pd.to_numeric, errors='coerce').fillna(0)
scaler = StandardScaler()
features = scaler.fit_transform(features)

#split & tensor conversion
X_train, X_val, y_train, y_val = train_test_split(features, targets, test_size=0.2,
X_train_tensor = torch.tensor(X_train, dtype=torch.float32)
y_train_tensor = torch.tensor(y_train, dtype=torch.float32)
X_val_tensor = torch.tensor(X_val, dtype=torch.float32)
y_val_tensor = torch.tensor(y_val, dtype=torch.float32)

#dataloader
train_dataset = TensorDataset(X_train_tensor, y_train_tensor)
val_dataset = TensorDataset(X_val_tensor, y_val_tensor)
train_loader = DataLoader(dataset=train_dataset, batch_size=32, shuffle=True)
val_loader = DataLoader(dataset=val_dataset, batch_size=32, shuffle=False)

#MLP model
class MLPHousing(nn.Module):
    def __init__(self, input_size):
        super(MLPHousing, self).__init__()
        self.fc1 = nn.Linear(input_size, 128)
        self.fc2 = nn.Linear(128, 64)
        self.fc3 = nn.Linear(64, 32)
        self.fc4 = nn.Linear(32, 1)

    def forward(self, x):
        x = torch.relu(self.fc1(x))
        x = torch.relu(self.fc2(x))
        x = torch.relu(self.fc3(x))
        x = self.fc4(x)
        return x

#model setup
input_size = X_train.shape[1]
model = MLPHousing(input_size)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model.to(device)
criterion = nn.MSELoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

#training Loop
epochs = 75
train_loss_history, val_loss_history = [], []
start_time = time.time()

for epoch in range(epochs):
    model.train()
    running_loss = 0.0
    for X_batch, y_batch in train_loader:
        X_batch, y_batch = X_batch.to(device), y_batch.to(device)
        optimizer.zero_grad()
        y_pred = model(X_batch)
        loss = criterion(y_pred, y_batch)
        loss.backward()
```

```
optimizer.step()
running_loss += loss.item() * X_batch.size(0)

train_loss = running_loss / len(train_loader.dataset)
train_loss_history.append(train_loss)

model.eval()
val_loss = 0.0
with torch.no_grad():
    for X_batch, y_batch in val_loader:
        X_batch, y_batch = X_batch.to(device), y_batch.to(device)
        y_pred = model(X_batch)
        loss = criterion(y_pred, y_batch)
        val_loss += loss.item() * X_batch.size(0)

    val_loss /= len(val_loader.dataset)
    val_loss_history.append(val_loss)

print(f"Epoch [{epoch+1}/{epochs}], Train Loss: {train_loss:.4f}, Val Loss: {va
end_time = time.time()
training_time = end_time - start_time

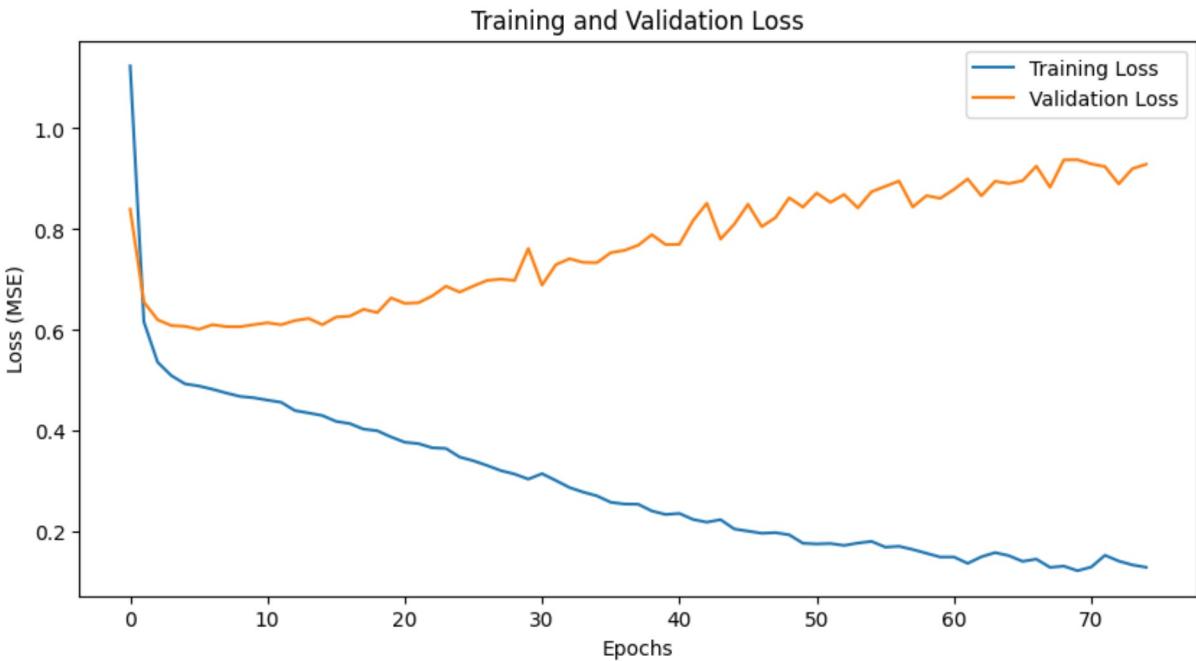
plt.figure(figsize=(10, 5))
plt.plot(train_loss_history, label='Training Loss')
plt.plot(val_loss_history, label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss (MSE)')
plt.title('Training and Validation Loss')
plt.legend()
plt.show()

model.eval()
with torch.no_grad():
    X_val_tensor, y_val_tensor = X_val_tensor.to(device), y_val_tensor.to(device)
    y_val_pred = model(X_val_tensor)
    mse = criterion(y_val_pred, y_val_tensor).item()
    rmse = np.sqrt(mse)

print(f"Final Validation MSE: {mse:.4f}")
print(f"Final Validation RMSE: {rmse:.4f}")
print(f"Model Complexity (Total Parameters): {sum(p.numel() for p in model.parameters())}")
print(f"Total Training Time: {training_time:.2f} seconds")
```

Epoch [1/75], Train Loss: 1.1236, Val Loss: 0.8398  
Epoch [2/75], Train Loss: 0.6164, Val Loss: 0.6556  
Epoch [3/75], Train Loss: 0.5364, Val Loss: 0.6206  
Epoch [4/75], Train Loss: 0.5098, Val Loss: 0.6091  
Epoch [5/75], Train Loss: 0.4934, Val Loss: 0.6075  
Epoch [6/75], Train Loss: 0.4892, Val Loss: 0.6017  
Epoch [7/75], Train Loss: 0.4828, Val Loss: 0.6107  
Epoch [8/75], Train Loss: 0.4753, Val Loss: 0.6069  
Epoch [9/75], Train Loss: 0.4686, Val Loss: 0.6067  
Epoch [10/75], Train Loss: 0.4659, Val Loss: 0.6108  
Epoch [11/75], Train Loss: 0.4612, Val Loss: 0.6146  
Epoch [12/75], Train Loss: 0.4569, Val Loss: 0.6107  
Epoch [13/75], Train Loss: 0.4405, Val Loss: 0.6189  
Epoch [14/75], Train Loss: 0.4358, Val Loss: 0.6232  
Epoch [15/75], Train Loss: 0.4307, Val Loss: 0.6107  
Epoch [16/75], Train Loss: 0.4191, Val Loss: 0.6259  
Epoch [17/75], Train Loss: 0.4146, Val Loss: 0.6277  
Epoch [18/75], Train Loss: 0.4037, Val Loss: 0.6412  
Epoch [19/75], Train Loss: 0.4005, Val Loss: 0.6347  
Epoch [20/75], Train Loss: 0.3882, Val Loss: 0.6639  
Epoch [21/75], Train Loss: 0.3777, Val Loss: 0.6531  
Epoch [22/75], Train Loss: 0.3750, Val Loss: 0.6543  
Epoch [23/75], Train Loss: 0.3667, Val Loss: 0.6678  
Epoch [24/75], Train Loss: 0.3654, Val Loss: 0.6871  
Epoch [25/75], Train Loss: 0.3483, Val Loss: 0.6755  
Epoch [26/75], Train Loss: 0.3410, Val Loss: 0.6874  
Epoch [27/75], Train Loss: 0.3316, Val Loss: 0.6985  
Epoch [28/75], Train Loss: 0.3214, Val Loss: 0.7011  
Epoch [29/75], Train Loss: 0.3146, Val Loss: 0.6985  
Epoch [30/75], Train Loss: 0.3046, Val Loss: 0.7619  
Epoch [31/75], Train Loss: 0.3154, Val Loss: 0.6893  
Epoch [32/75], Train Loss: 0.3020, Val Loss: 0.7296  
Epoch [33/75], Train Loss: 0.2879, Val Loss: 0.7415  
Epoch [34/75], Train Loss: 0.2788, Val Loss: 0.7343  
Epoch [35/75], Train Loss: 0.2714, Val Loss: 0.7338  
Epoch [36/75], Train Loss: 0.2587, Val Loss: 0.7535  
Epoch [37/75], Train Loss: 0.2553, Val Loss: 0.7581  
Epoch [38/75], Train Loss: 0.2548, Val Loss: 0.7683  
Epoch [39/75], Train Loss: 0.2414, Val Loss: 0.7893  
Epoch [40/75], Train Loss: 0.2346, Val Loss: 0.7695  
Epoch [41/75], Train Loss: 0.2366, Val Loss: 0.7700  
Epoch [42/75], Train Loss: 0.2246, Val Loss: 0.8172  
Epoch [43/75], Train Loss: 0.2192, Val Loss: 0.8515  
Epoch [44/75], Train Loss: 0.2243, Val Loss: 0.7804  
Epoch [45/75], Train Loss: 0.2058, Val Loss: 0.8100  
Epoch [46/75], Train Loss: 0.2016, Val Loss: 0.8496  
Epoch [47/75], Train Loss: 0.1973, Val Loss: 0.8052  
Epoch [48/75], Train Loss: 0.1985, Val Loss: 0.8230  
Epoch [49/75], Train Loss: 0.1942, Val Loss: 0.8624  
Epoch [50/75], Train Loss: 0.1775, Val Loss: 0.8439  
Epoch [51/75], Train Loss: 0.1761, Val Loss: 0.8718  
Epoch [52/75], Train Loss: 0.1770, Val Loss: 0.8534  
Epoch [53/75], Train Loss: 0.1732, Val Loss: 0.8688  
Epoch [54/75], Train Loss: 0.1778, Val Loss: 0.8423  
Epoch [55/75], Train Loss: 0.1811, Val Loss: 0.8744  
Epoch [56/75], Train Loss: 0.1698, Val Loss: 0.8850

```
Epoch [57/75], Train Loss: 0.1714, Val Loss: 0.8955
Epoch [58/75], Train Loss: 0.1651, Val Loss: 0.8441
Epoch [59/75], Train Loss: 0.1576, Val Loss: 0.8666
Epoch [60/75], Train Loss: 0.1502, Val Loss: 0.8613
Epoch [61/75], Train Loss: 0.1502, Val Loss: 0.8790
Epoch [62/75], Train Loss: 0.1376, Val Loss: 0.8997
Epoch [63/75], Train Loss: 0.1507, Val Loss: 0.8663
Epoch [64/75], Train Loss: 0.1589, Val Loss: 0.8952
Epoch [65/75], Train Loss: 0.1528, Val Loss: 0.8905
Epoch [66/75], Train Loss: 0.1416, Val Loss: 0.8962
Epoch [67/75], Train Loss: 0.1460, Val Loss: 0.9251
Epoch [68/75], Train Loss: 0.1296, Val Loss: 0.8832
Epoch [69/75], Train Loss: 0.1321, Val Loss: 0.9375
Epoch [70/75], Train Loss: 0.1231, Val Loss: 0.9378
Epoch [71/75], Train Loss: 0.1303, Val Loss: 0.9295
Epoch [72/75], Train Loss: 0.1537, Val Loss: 0.9242
Epoch [73/75], Train Loss: 0.1421, Val Loss: 0.8898
Epoch [74/75], Train Loss: 0.1345, Val Loss: 0.9201
Epoch [75/75], Train Loss: 0.1298, Val Loss: 0.9288
```



```
Final Validation MSE: 0.9288
Final Validation RMSE: 0.9637
Model Complexity (Total Parameters): 12033
Total Training Time: 1.77 seconds
```

```
In [12]: # Import required libraries
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import time
```

```
# Load and prepare dataset
data_source = "https://raw.githubusercontent.com/HamedTabkhi/Intro-to-ML/refs/heads/main.csv"
house_data = pd.read_csv(data_source)

# Convert categorical variables to numeric
for column in house_data.select_dtypes(include=['object']).columns:
    house_data[column] = house_data[column].astype('category').cat.codes

# Separate features and target
price_column = house_data.columns[-1]
input_features = house_data.drop([price_column], axis=1)
house_prices = house_data[price_column].values.reshape(-1, 1)

# Clean and scale the data
input_features = input_features.apply(pd.to_numeric, errors='coerce').fillna(0)
feature_scaler = StandardScaler()
input_features = feature_scaler.fit_transform(input_features)

# Split data into training and validation sets
X_train, X_valid, y_train, y_valid = train_test_split(input_features, house_prices,
X_train_tensor = torch.tensor(X_train, dtype=torch.float32)
y_train_tensor = torch.tensor(y_train, dtype=torch.float32)
X_valid_tensor = torch.tensor(X_valid, dtype=torch.float32)
y_valid_tensor = torch.tensor(y_valid, dtype=torch.float32)

# Create data loaders
train_dataset = TensorDataset(X_train_tensor, y_train_tensor)
valid_dataset = TensorDataset(X_valid_tensor, y_valid_tensor)
train_loader = DataLoader(dataset=train_dataset, batch_size=32, shuffle=True)
valid_loader = DataLoader(dataset=valid_dataset, batch_size=32, shuffle=False)

# Define neural network architecture
class HousePriceNet(nn.Module):
    def __init__(self, input_dim):
        super(HousePriceNet, self).__init__()
        self.layer1 = nn.Linear(input_dim, 128)
        self.layer2 = nn.Linear(128, 64)
        self.layer3 = nn.Linear(64, 32)
        self.layer4 = nn.Linear(32, 1)

    def forward(self, x):
        x = torch.relu(self.layer1(x))
        x = torch.relu(self.layer2(x))
        x = torch.relu(self.layer3(x))
        x = self.layer4(x)
        return x

# Initialize model and training components
input_dim = X_train.shape[1]
model = HousePriceNet(input_dim)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model.to(device)
loss_function = nn.MSELoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Training Loop
```

```
num_epochs = 75
train_losses, valid_losses = [], []
start_time = time.time()

for epoch in range(num_epochs):
    model.train()
    epoch_loss = 0.0
    for batch_X, batch_y in train_loader:
        batch_X, batch_y = batch_X.to(device), batch_y.to(device)
        optimizer.zero_grad()
        predictions = model(batch_X)
        loss = loss_function(predictions, batch_y)
        loss.backward()
        optimizer.step()
        epoch_loss += loss.item() * batch_X.size(0)

    avg_train_loss = epoch_loss / len(train_loader.dataset)
    train_losses.append(avg_train_loss)

    model.eval()
    valid_epoch_loss = 0.0
    with torch.no_grad():
        for batch_X, batch_y in valid_loader:
            batch_X, batch_y = batch_X.to(device), batch_y.to(device)
            predictions = model(batch_X)
            loss = loss_function(predictions, batch_y)
            valid_epoch_loss += loss.item() * batch_X.size(0)

    avg_valid_loss = valid_epoch_loss / len(valid_loader.dataset)
    valid_losses.append(avg_valid_loss)

print(f"Epoch [{epoch+1}/{num_epochs}], Training Loss: {avg_train_loss:.4f}, Va
total_time = time.time() - start_time

# Plot training results
plt.figure(figsize=(10, 5))
plt.plot(train_losses, label='Training Loss')
plt.plot(valid_losses, label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss (MSE)')
plt.title('Training and Validation Loss Over Time')
plt.legend()
plt.show()

# Final evaluation
model.eval()
with torch.no_grad():
    X_valid_tensor, y_valid_tensor = X_valid_tensor.to(device), y_valid_tensor.to(d
    final_predictions = model(X_valid_tensor)
    final_mse = loss_function(final_predictions, y_valid_tensor).item()
    final_rmse = np.sqrt(final_mse)

print(f"Final MSE on Validation Set: {final_mse:.4f}")
print(f"Final RMSE on Validation Set: {final_rmse:.4f}")
print(f"Total Parameters in Model: {sum(p.numel() for p in model.parameters())}"
```

```
print(f"Training Duration: {total_time:.2f} seconds")
```

Epoch [1/75], Training Loss: 1.0700, Validation Loss: 0.8127  
Epoch [2/75], Training Loss: 0.6071, Validation Loss: 0.6170  
Epoch [3/75], Training Loss: 0.5154, Validation Loss: 0.6109  
Epoch [4/75], Training Loss: 0.5041, Validation Loss: 0.6056  
Epoch [5/75], Training Loss: 0.4967, Validation Loss: 0.6134  
Epoch [6/75], Training Loss: 0.4867, Validation Loss: 0.6021  
Epoch [7/75], Training Loss: 0.4733, Validation Loss: 0.6004  
Epoch [8/75], Training Loss: 0.4672, Validation Loss: 0.5995  
Epoch [9/75], Training Loss: 0.4644, Validation Loss: 0.6101  
Epoch [10/75], Training Loss: 0.4515, Validation Loss: 0.6062  
Epoch [11/75], Training Loss: 0.4452, Validation Loss: 0.6068  
Epoch [12/75], Training Loss: 0.4337, Validation Loss: 0.6082  
Epoch [13/75], Training Loss: 0.4276, Validation Loss: 0.6175  
Epoch [14/75], Training Loss: 0.4200, Validation Loss: 0.6301  
Epoch [15/75], Training Loss: 0.4112, Validation Loss: 0.6336  
Epoch [16/75], Training Loss: 0.4056, Validation Loss: 0.6316  
Epoch [17/75], Training Loss: 0.3942, Validation Loss: 0.6388  
Epoch [18/75], Training Loss: 0.3874, Validation Loss: 0.6423  
Epoch [19/75], Training Loss: 0.3803, Validation Loss: 0.6638  
Epoch [20/75], Training Loss: 0.3672, Validation Loss: 0.6603  
Epoch [21/75], Training Loss: 0.3617, Validation Loss: 0.6611  
Epoch [22/75], Training Loss: 0.3478, Validation Loss: 0.6811  
Epoch [23/75], Training Loss: 0.3326, Validation Loss: 0.6748  
Epoch [24/75], Training Loss: 0.3243, Validation Loss: 0.6759  
Epoch [25/75], Training Loss: 0.3213, Validation Loss: 0.7173  
Epoch [26/75], Training Loss: 0.3157, Validation Loss: 0.6890  
Epoch [27/75], Training Loss: 0.3174, Validation Loss: 0.6957  
Epoch [28/75], Training Loss: 0.2922, Validation Loss: 0.7241  
Epoch [29/75], Training Loss: 0.3101, Validation Loss: 0.7307  
Epoch [30/75], Training Loss: 0.2808, Validation Loss: 0.7111  
Epoch [31/75], Training Loss: 0.2771, Validation Loss: 0.7301  
Epoch [32/75], Training Loss: 0.2721, Validation Loss: 0.7403  
Epoch [33/75], Training Loss: 0.2591, Validation Loss: 0.7334  
Epoch [34/75], Training Loss: 0.2486, Validation Loss: 0.7248  
Epoch [35/75], Training Loss: 0.2488, Validation Loss: 0.7624  
Epoch [36/75], Training Loss: 0.2441, Validation Loss: 0.7552  
Epoch [37/75], Training Loss: 0.2379, Validation Loss: 0.7818  
Epoch [38/75], Training Loss: 0.2420, Validation Loss: 0.7634  
Epoch [39/75], Training Loss: 0.2305, Validation Loss: 0.7816  
Epoch [40/75], Training Loss: 0.2201, Validation Loss: 0.7642  
Epoch [41/75], Training Loss: 0.2201, Validation Loss: 0.8012  
Epoch [42/75], Training Loss: 0.2248, Validation Loss: 0.7749  
Epoch [43/75], Training Loss: 0.2141, Validation Loss: 0.8332  
Epoch [44/75], Training Loss: 0.2047, Validation Loss: 0.7821  
Epoch [45/75], Training Loss: 0.1956, Validation Loss: 0.8155  
Epoch [46/75], Training Loss: 0.2009, Validation Loss: 0.8855  
Epoch [47/75], Training Loss: 0.2010, Validation Loss: 0.8590  
Epoch [48/75], Training Loss: 0.1944, Validation Loss: 0.8717  
Epoch [49/75], Training Loss: 0.1953, Validation Loss: 0.8179  
Epoch [50/75], Training Loss: 0.1755, Validation Loss: 0.8570  
Epoch [51/75], Training Loss: 0.1735, Validation Loss: 0.8712  
Epoch [52/75], Training Loss: 0.1780, Validation Loss: 0.8881  
Epoch [53/75], Training Loss: 0.1705, Validation Loss: 0.8757  
Epoch [54/75], Training Loss: 0.1552, Validation Loss: 0.8854  
Epoch [55/75], Training Loss: 0.1542, Validation Loss: 0.9129  
Epoch [56/75], Training Loss: 0.1617, Validation Loss: 0.9158

```
Epoch [57/75], Training Loss: 0.1642, Validation Loss: 0.8993
Epoch [58/75], Training Loss: 0.1636, Validation Loss: 0.9402
Epoch [59/75], Training Loss: 0.1674, Validation Loss: 0.8704
Epoch [60/75], Training Loss: 0.1676, Validation Loss: 0.9278
Epoch [61/75], Training Loss: 0.1531, Validation Loss: 0.9271
Epoch [62/75], Training Loss: 0.1393, Validation Loss: 0.9258
Epoch [63/75], Training Loss: 0.1376, Validation Loss: 0.9676
Epoch [64/75], Training Loss: 0.1348, Validation Loss: 0.9617
Epoch [65/75], Training Loss: 0.1286, Validation Loss: 0.9681
Epoch [66/75], Training Loss: 0.1293, Validation Loss: 0.9794
Epoch [67/75], Training Loss: 0.1255, Validation Loss: 0.9728
Epoch [68/75], Training Loss: 0.1225, Validation Loss: 1.0135
Epoch [69/75], Training Loss: 0.1266, Validation Loss: 0.9988
Epoch [70/75], Training Loss: 0.1202, Validation Loss: 0.9914
Epoch [71/75], Training Loss: 0.1223, Validation Loss: 1.0041
Epoch [72/75], Training Loss: 0.1332, Validation Loss: 0.9811
Epoch [73/75], Training Loss: 0.1334, Validation Loss: 0.9820
Epoch [74/75], Training Loss: 0.1311, Validation Loss: 1.0073
Epoch [75/75], Training Loss: 0.1343, Validation Loss: 1.0022
```



Final MSE on Validation Set: 1.0022

Final RMSE on Validation Set: 1.0011

Total Parameters in Model: 12033

Training Duration: 1.93 seconds

```
In [11]: #2a
# Import required libraries
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import time
```

```
# Load and preprocess the housing dataset
url = "https://raw.githubusercontent.com/HamedTabkhi/Intro-to-ML/refs/heads/main/Da
data = pd.read_csv(url)

# Convert categorical variables to numeric
for col in data.select_dtypes(include=['object']).columns:
    data[col] = data[col].astype('category').cat.codes

# Separate features and target
target_column = data.columns[-1] # 'price' is the last column
features = data.drop([target_column], axis=1)
targets = data[target_column].values.reshape(-1, 1)

# Handle any missing values and convert to numeric
features = features.apply(pd.to_numeric, errors='coerce').fillna(0)

# Scale the features
scaler = StandardScaler()
features = scaler.fit_transform(features)

# Split data and convert to PyTorch tensors
X_train, X_val, y_train, y_val = train_test_split(features, targets, test_size=0.2,
X_train_tensor = torch.tensor(X_train, dtype=torch.float32)
y_train_tensor = torch.tensor(y_train, dtype=torch.float32)
X_val_tensor = torch.tensor(X_val, dtype=torch.float32)
y_val_tensor = torch.tensor(y_val, dtype=torch.float32)

# Create DataLoader objects for batch processing
train_dataset = TensorDataset(X_train_tensor, y_train_tensor)
val_dataset = TensorDataset(X_val_tensor, y_val_tensor)
train_loader = DataLoader(dataset=train_dataset, batch_size=32, shuffle=True)
val_loader = DataLoader(dataset=val_dataset, batch_size=32, shuffle=False)

# Define the MLP model for housing price prediction
class MLPHousing(nn.Module):
    def __init__(self, input_size):
        super(MLPHousing, self).__init__()
        # Modified architecture with gradually decreasing layer sizes
        self.model = nn.Sequential(
            nn.Linear(input_size, 256),
            nn.ReLU(),
            nn.Linear(256, 128),
            nn.ReLU(),
            nn.Linear(128, 64),
            nn.ReLU(),
            nn.Linear(64, 1)
        )

        def forward(self, x):
            return self.model(x)

# Initialize model, loss function, and optimizer
input_size = X_train.shape[1]
model = MLPHousing(input_size)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

```
model.to(device)
criterion = nn.MSELoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Training Loop
epochs = 75
train_loss_history = []
val_loss_history = []
start_time = time.time()

for epoch in range(epochs):
    # Training phase
    model.train()
    running_loss = 0.0
    for X_batch, y_batch in train_loader:
        X_batch, y_batch = X_batch.to(device), y_batch.to(device)

        # Forward pass and loss calculation
        optimizer.zero_grad()
        y_pred = model(X_batch)
        loss = criterion(y_pred, y_batch)

        # Backward pass and optimization
        loss.backward()
        optimizer.step()

        running_loss += loss.item() * X_batch.size(0)

    train_loss = running_loss / len(train_loader.dataset)
    train_loss_history.append(train_loss)

    # Validation phase
    model.eval()
    val_loss = 0.0
    val_predictions = []
    val_targets = []

    with torch.no_grad():
        for X_batch, y_batch in val_loader:
            X_batch, y_batch = X_batch.to(device), y_batch.to(device)
            y_pred = model(X_batch)
            loss = criterion(y_pred, y_batch)
            val_loss += loss.item() * X_batch.size(0)

            # Store predictions and targets for additional metrics
            val_predictions.extend(y_pred.cpu().numpy())
            val_targets.extend(y_batch.cpu().numpy())

    val_loss /= len(val_loader.dataset)
    val_loss_history.append(val_loss)

    print(f"Epoch [{epoch+1}/{epochs}], Train Loss: {train_loss:.4f}, Val Loss: {va
end_time = time.time()
training_time = end_time - start_time
```

```
# Plot training results
plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(train_loss_history, label='Training Loss')
plt.plot(val_loss_history, label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss (MSE)')
plt.title('Training and Validation Loss')
plt.legend()

# Plot predictions vs actual values
plt.subplot(1, 2, 2)
val_predictions = np.array(val_predictions)
val_targets = np.array(val_targets)
plt.scatter(val_targets, val_predictions, alpha=0.5)
plt.plot([val_targets.min(), val_targets.max()],
         [val_targets.min(), val_targets.max()],
         'r--', label='Perfect Prediction')
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.title('Predictions vs Actual Values')
plt.legend()
plt.tight_layout()
plt.show()

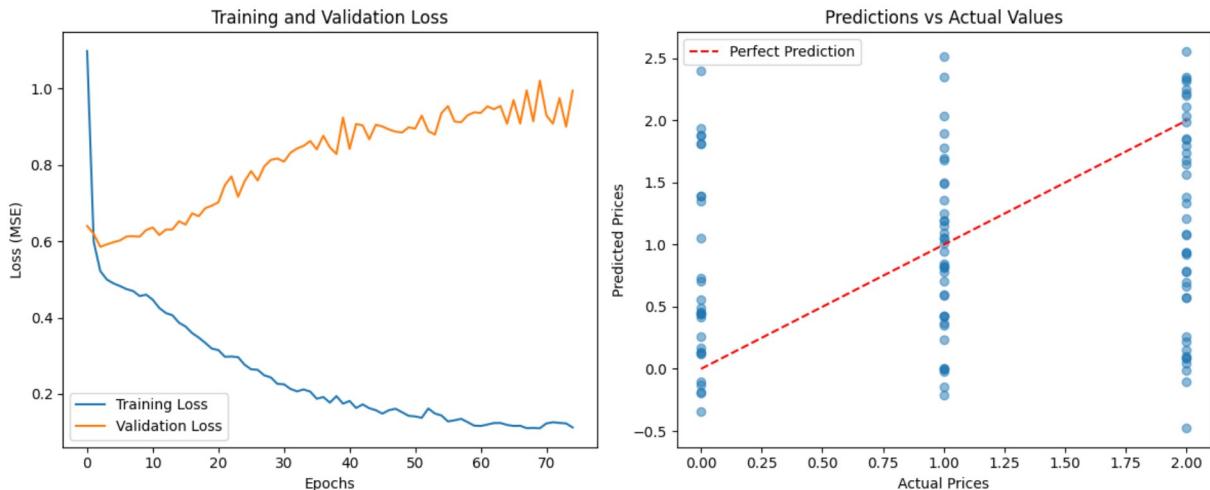
# Calculate and display final metrics
model.eval()
with torch.no_grad():
    X_val_tensor, y_val_tensor = X_val_tensor.to(device), y_val_tensor.to(device)
    y_val_pred = model(X_val_tensor)
    mse = criterion(y_val_pred, y_val_tensor).item()
    rmse = np.sqrt(mse)

    # Calculate R2 score
    y_val_mean = torch.mean(y_val_tensor)
    ss_tot = torch.sum((y_val_tensor - y_val_mean) ** 2)
    ss_res = torch.sum((y_val_tensor - y_val_pred) ** 2)
    r2 = 1 - (ss_res / ss_tot)

print("\nFinal Model Performance:")
print(f"Validation MSE: {mse:.4f}")
print(f"Validation RMSE: {rmse:.4f}")
print(f"R2 Score: {r2.item():.4f}")
print(f"Model Parameters: {sum(p.numel() for p in model.parameters())}")
print(f"Training Time: {training_time:.2f} seconds")
```

Epoch [1/75], Train Loss: 1.0983, Val Loss: 0.6401  
Epoch [2/75], Train Loss: 0.5991, Val Loss: 0.6196  
Epoch [3/75], Train Loss: 0.5221, Val Loss: 0.5854  
Epoch [4/75], Train Loss: 0.4997, Val Loss: 0.5919  
Epoch [5/75], Train Loss: 0.4898, Val Loss: 0.5974  
Epoch [6/75], Train Loss: 0.4827, Val Loss: 0.6021  
Epoch [7/75], Train Loss: 0.4746, Val Loss: 0.6125  
Epoch [8/75], Train Loss: 0.4697, Val Loss: 0.6132  
Epoch [9/75], Train Loss: 0.4565, Val Loss: 0.6122  
Epoch [10/75], Train Loss: 0.4601, Val Loss: 0.6295  
Epoch [11/75], Train Loss: 0.4468, Val Loss: 0.6362  
Epoch [12/75], Train Loss: 0.4252, Val Loss: 0.6164  
Epoch [13/75], Train Loss: 0.4120, Val Loss: 0.6306  
Epoch [14/75], Train Loss: 0.4061, Val Loss: 0.6310  
Epoch [15/75], Train Loss: 0.3866, Val Loss: 0.6524  
Epoch [16/75], Train Loss: 0.3766, Val Loss: 0.6437  
Epoch [17/75], Train Loss: 0.3595, Val Loss: 0.6730  
Epoch [18/75], Train Loss: 0.3478, Val Loss: 0.6656  
Epoch [19/75], Train Loss: 0.3339, Val Loss: 0.6864  
Epoch [20/75], Train Loss: 0.3189, Val Loss: 0.6932  
Epoch [21/75], Train Loss: 0.3149, Val Loss: 0.7021  
Epoch [22/75], Train Loss: 0.2973, Val Loss: 0.7474  
Epoch [23/75], Train Loss: 0.2980, Val Loss: 0.7698  
Epoch [24/75], Train Loss: 0.2965, Val Loss: 0.7164  
Epoch [25/75], Train Loss: 0.2766, Val Loss: 0.7572  
Epoch [26/75], Train Loss: 0.2645, Val Loss: 0.7838  
Epoch [27/75], Train Loss: 0.2631, Val Loss: 0.7592  
Epoch [28/75], Train Loss: 0.2486, Val Loss: 0.7963  
Epoch [29/75], Train Loss: 0.2433, Val Loss: 0.8132  
Epoch [30/75], Train Loss: 0.2265, Val Loss: 0.8168  
Epoch [31/75], Train Loss: 0.2253, Val Loss: 0.8086  
Epoch [32/75], Train Loss: 0.2131, Val Loss: 0.8319  
Epoch [33/75], Train Loss: 0.2068, Val Loss: 0.8431  
Epoch [34/75], Train Loss: 0.2117, Val Loss: 0.8500  
Epoch [35/75], Train Loss: 0.2062, Val Loss: 0.8626  
Epoch [36/75], Train Loss: 0.1874, Val Loss: 0.8409  
Epoch [37/75], Train Loss: 0.1919, Val Loss: 0.8764  
Epoch [38/75], Train Loss: 0.1772, Val Loss: 0.8459  
Epoch [39/75], Train Loss: 0.1943, Val Loss: 0.8286  
Epoch [40/75], Train Loss: 0.1745, Val Loss: 0.9243  
Epoch [41/75], Train Loss: 0.1816, Val Loss: 0.8421  
Epoch [42/75], Train Loss: 0.1632, Val Loss: 0.9073  
Epoch [43/75], Train Loss: 0.1726, Val Loss: 0.9034  
Epoch [44/75], Train Loss: 0.1627, Val Loss: 0.8673  
Epoch [45/75], Train Loss: 0.1575, Val Loss: 0.9053  
Epoch [46/75], Train Loss: 0.1484, Val Loss: 0.9007  
Epoch [47/75], Train Loss: 0.1575, Val Loss: 0.8929  
Epoch [48/75], Train Loss: 0.1613, Val Loss: 0.8871  
Epoch [49/75], Train Loss: 0.1521, Val Loss: 0.8850  
Epoch [50/75], Train Loss: 0.1426, Val Loss: 0.8985  
Epoch [51/75], Train Loss: 0.1412, Val Loss: 0.8951  
Epoch [52/75], Train Loss: 0.1374, Val Loss: 0.9292  
Epoch [53/75], Train Loss: 0.1619, Val Loss: 0.8887  
Epoch [54/75], Train Loss: 0.1488, Val Loss: 0.8795  
Epoch [55/75], Train Loss: 0.1437, Val Loss: 0.9355  
Epoch [56/75], Train Loss: 0.1279, Val Loss: 0.9541

```
Epoch [57/75], Train Loss: 0.1310, Val Loss: 0.9138
Epoch [58/75], Train Loss: 0.1348, Val Loss: 0.9119
Epoch [59/75], Train Loss: 0.1261, Val Loss: 0.9303
Epoch [60/75], Train Loss: 0.1170, Val Loss: 0.9379
Epoch [61/75], Train Loss: 0.1161, Val Loss: 0.9362
Epoch [62/75], Train Loss: 0.1198, Val Loss: 0.9533
Epoch [63/75], Train Loss: 0.1238, Val Loss: 0.9461
Epoch [64/75], Train Loss: 0.1240, Val Loss: 0.9543
Epoch [65/75], Train Loss: 0.1189, Val Loss: 0.9078
Epoch [66/75], Train Loss: 0.1162, Val Loss: 0.9699
Epoch [67/75], Train Loss: 0.1166, Val Loss: 0.9084
Epoch [68/75], Train Loss: 0.1104, Val Loss: 0.9952
Epoch [69/75], Train Loss: 0.1109, Val Loss: 0.9142
Epoch [70/75], Train Loss: 0.1101, Val Loss: 1.0209
Epoch [71/75], Train Loss: 0.1228, Val Loss: 0.9299
Epoch [72/75], Train Loss: 0.1258, Val Loss: 0.9079
Epoch [73/75], Train Loss: 0.1241, Val Loss: 0.9751
Epoch [74/75], Train Loss: 0.1228, Val Loss: 0.9000
Epoch [75/75], Train Loss: 0.1122, Val Loss: 0.9943
```



#### Final Model Performance:

```
Validation MSE: 0.9943
Validation RMSE: 0.9972
R2 Score: -0.5781
Model Parameters: 44545
Training Time: 1.79 seconds
```

```
In [ ]: #2b
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import time

# Load housing dataset from GitHub repository
data_url = "https://raw.githubusercontent.com/HamedTabkhi/Intro-to-ML/refs/heads/main/housing_dataset.csv"
housing_dataset = pd.read_csv(data_url)
```

```
# Convert categorical variables to numeric codes
for col in housing_dataset.select_dtypes(include=['object']).columns:
    housing_dataset[col] = housing_dataset[col].astype('category').cat.codes

# Perform one-hot encoding on categorical variables
feature_encoder = OneHotEncoder(sparse_output=False, handle_unknown='ignore')
categorical_cols = housing_dataset.select_dtypes(include=['category', 'object']).columns
encoded_features = feature_encoder.fit_transform(housing_dataset[categorical_cols])
encoded_features = pd.DataFrame(encoded_features, columns=feature_encoder.get_feature_names_out())
housing_dataset = pd.concat([housing_dataset.drop(categorical_cols, axis=1), encoded_features], axis=1)

# Define target variable and separate features
target_column = "SalePrice"
if target_column not in housing_dataset.columns:
    target_column = housing_dataset.columns[-1]

input_data = housing_dataset.drop([target_column], axis=1)
price_targets = housing_dataset[target_column].values.reshape(-1, 1)

# Clean and scale the input features
input_data = input_data.apply(pd.to_numeric, errors='coerce').fillna(0)
data_scaler = StandardScaler()
input_data = data_scaler.fit_transform(input_data)

# Split data into training and validation sets
train_features, val_features, train_targets, val_targets = train_test_split(
    input_data, price_targets, test_size=0.2, random_state=42
)

# Convert numpy arrays to PyTorch tensors
train_features_tensor = torch.tensor(train_features, dtype=torch.float32)
train_targets_tensor = torch.tensor(train_targets, dtype=torch.float32)
val_features_tensor = torch.tensor(val_features, dtype=torch.float32)
val_targets_tensor = torch.tensor(val_targets, dtype=torch.float32)

# Create DataLoader objects for batch processing
train_dataset = TensorDataset(train_features_tensor, train_targets_tensor)
val_dataset = TensorDataset(val_features_tensor, val_targets_tensor)
train_loader = DataLoader(dataset=train_dataset, batch_size=32, shuffle=True)
val_loader = DataLoader(dataset=val_dataset, batch_size=32, shuffle=False)

# Define neural network architecture for house price prediction
class HousePriceModel(nn.Module):
    def __init__(self, input_size):
        super(HousePriceModel, self).__init__()
        # Define network layers with descriptive names
        self.input_layer = nn.Linear(input_size, 128)
        self.hidden_layer1 = nn.Linear(128, 64)
        self.hidden_layer2 = nn.Linear(64, 32)
        self.output_layer = nn.Linear(32, 1)

    def forward(self, x):
        # Apply ReLU activation after each layer except the output
        x = torch.relu(self.input_layer(x))
        x = torch.relu(self.hidden_layer1(x))
```

```
x = torch.relu(self.hidden_layer2(x))
x = self.output_layer(x)
return x

# Initialize model and training components
input_dim = train_features.shape[1]
model = HousePriceModel(input_dim)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model.to(device)
loss_function = nn.MSELoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Training settings
num_epochs = 75
training_losses = []
validation_losses = []
start_time = time.time()

# Training Loop
for epoch in range(num_epochs):
    # Training phase
    model.train()
    batch_loss = 0.0
    for batch_features, batch_targets in train_loader:
        # Move batch data to appropriate device
        batch_features = batch_features.to(device)
        batch_targets = batch_targets.to(device)

        # Forward pass and loss calculation
        optimizer.zero_grad()
        predictions = model(batch_features)
        loss = loss_function(predictions, batch_targets)

        # Backward pass and optimization
        loss.backward()
        optimizer.step()
        batch_loss += loss.item() * batch_features.size(0)

    # Calculate average training loss
    avg_train_loss = batch_loss / len(train_loader.dataset)
    training_losses.append(avg_train_loss)

    # Validation phase
    model.eval()
    val_batch_loss = 0.0
    with torch.no_grad():
        for batch_features, batch_targets in val_loader:
            batch_features = batch_features.to(device)
            batch_targets = batch_targets.to(device)
            predictions = model(batch_features)
            loss = loss_function(predictions, batch_targets)
            val_batch_loss += loss.item() * batch_features.size(0)

    # Calculate average validation loss
    avg_val_loss = val_batch_loss / len(val_loader.dataset)
    validation_losses.append(avg_val_loss)
```

```
# Print progress
print(f"Epoch [{epoch+1}/{num_epochs}], Training Loss: {avg_train_loss:.4f}, Va

# Calculate total training time
total_time = time.time() - start_time

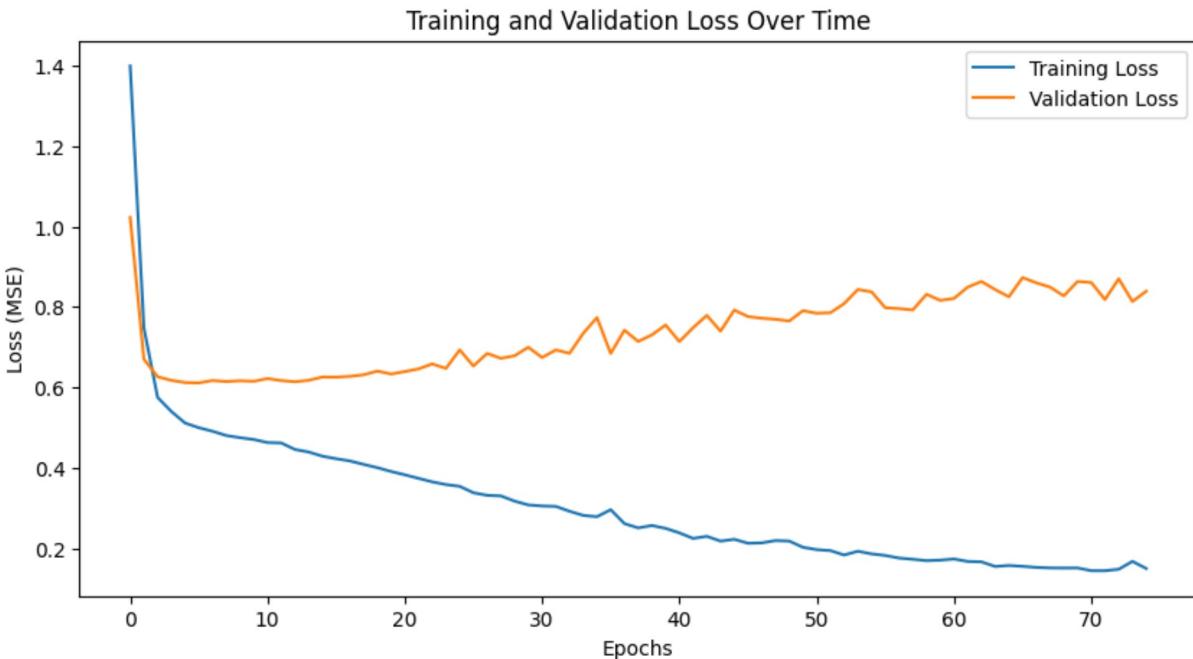
# Visualize training progress
plt.figure(figsize=(10, 5))
plt.plot(training_losses, label='Training Loss')
plt.plot(validation_losses, label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss (MSE)')
plt.title('Training and Validation Loss Over Time')
plt.legend()
plt.show()

# Final model evaluation
model.eval()
with torch.no_grad():
    val_features_tensor = val_features_tensor.to(device)
    val_targets_tensor = val_targets_tensor.to(device)
    final_predictions = model(val_features_tensor)
    final_mse = loss_function(final_predictions, val_targets_tensor).item()
    final_rmse = np.sqrt(final_mse)

# Print final metrics
print(f"Final Validation MSE: {final_mse:.4f}")
print(f"Final Validation RMSE: {final_rmse:.4f}")
print(f"Model Complexity (Total Parameters): {sum(p.numel() for p in model.parameters())}")
print(f"Total Training Time: {total_time:.2f} seconds")
```

Epoch [1/75], Training Loss: 1.3997, Validation Loss: 1.0229  
Epoch [2/75], Training Loss: 0.7478, Validation Loss: 0.6702  
Epoch [3/75], Training Loss: 0.5750, Validation Loss: 0.6264  
Epoch [4/75], Training Loss: 0.5400, Validation Loss: 0.6174  
Epoch [5/75], Training Loss: 0.5112, Validation Loss: 0.6120  
Epoch [6/75], Training Loss: 0.4996, Validation Loss: 0.6111  
Epoch [7/75], Training Loss: 0.4908, Validation Loss: 0.6170  
Epoch [8/75], Training Loss: 0.4801, Validation Loss: 0.6145  
Epoch [9/75], Training Loss: 0.4748, Validation Loss: 0.6164  
Epoch [10/75], Training Loss: 0.4702, Validation Loss: 0.6151  
Epoch [11/75], Training Loss: 0.4626, Validation Loss: 0.6218  
Epoch [12/75], Training Loss: 0.4617, Validation Loss: 0.6169  
Epoch [13/75], Training Loss: 0.4455, Validation Loss: 0.6139  
Epoch [14/75], Training Loss: 0.4391, Validation Loss: 0.6176  
Epoch [15/75], Training Loss: 0.4288, Validation Loss: 0.6257  
Epoch [16/75], Training Loss: 0.4225, Validation Loss: 0.6254  
Epoch [17/75], Training Loss: 0.4169, Validation Loss: 0.6274  
Epoch [18/75], Training Loss: 0.4083, Validation Loss: 0.6312  
Epoch [19/75], Training Loss: 0.4002, Validation Loss: 0.6404  
Epoch [20/75], Training Loss: 0.3907, Validation Loss: 0.6332  
Epoch [21/75], Training Loss: 0.3824, Validation Loss: 0.6392  
Epoch [22/75], Training Loss: 0.3736, Validation Loss: 0.6459  
Epoch [23/75], Training Loss: 0.3648, Validation Loss: 0.6582  
Epoch [24/75], Training Loss: 0.3581, Validation Loss: 0.6471  
Epoch [25/75], Training Loss: 0.3537, Validation Loss: 0.6933  
Epoch [26/75], Training Loss: 0.3377, Validation Loss: 0.6531  
Epoch [27/75], Training Loss: 0.3313, Validation Loss: 0.6844  
Epoch [28/75], Training Loss: 0.3299, Validation Loss: 0.6723  
Epoch [29/75], Training Loss: 0.3169, Validation Loss: 0.6785  
Epoch [30/75], Training Loss: 0.3074, Validation Loss: 0.6998  
Epoch [31/75], Training Loss: 0.3048, Validation Loss: 0.6742  
Epoch [32/75], Training Loss: 0.3038, Validation Loss: 0.6931  
Epoch [33/75], Training Loss: 0.2917, Validation Loss: 0.6847  
Epoch [34/75], Training Loss: 0.2814, Validation Loss: 0.7345  
Epoch [35/75], Training Loss: 0.2781, Validation Loss: 0.7737  
Epoch [36/75], Training Loss: 0.2956, Validation Loss: 0.6843  
Epoch [37/75], Training Loss: 0.2609, Validation Loss: 0.7423  
Epoch [38/75], Training Loss: 0.2503, Validation Loss: 0.7142  
Epoch [39/75], Training Loss: 0.2562, Validation Loss: 0.7305  
Epoch [40/75], Training Loss: 0.2491, Validation Loss: 0.7551  
Epoch [41/75], Training Loss: 0.2379, Validation Loss: 0.7140  
Epoch [42/75], Training Loss: 0.2241, Validation Loss: 0.7484  
Epoch [43/75], Training Loss: 0.2293, Validation Loss: 0.7791  
Epoch [44/75], Training Loss: 0.2174, Validation Loss: 0.7398  
Epoch [45/75], Training Loss: 0.2219, Validation Loss: 0.7924  
Epoch [46/75], Training Loss: 0.2121, Validation Loss: 0.7761  
Epoch [47/75], Training Loss: 0.2130, Validation Loss: 0.7719  
Epoch [48/75], Training Loss: 0.2188, Validation Loss: 0.7695  
Epoch [49/75], Training Loss: 0.2175, Validation Loss: 0.7651  
Epoch [50/75], Training Loss: 0.2020, Validation Loss: 0.7907  
Epoch [51/75], Training Loss: 0.1962, Validation Loss: 0.7844  
Epoch [52/75], Training Loss: 0.1938, Validation Loss: 0.7856  
Epoch [53/75], Training Loss: 0.1829, Validation Loss: 0.8084  
Epoch [54/75], Training Loss: 0.1921, Validation Loss: 0.8435  
Epoch [55/75], Training Loss: 0.1856, Validation Loss: 0.8376  
Epoch [56/75], Training Loss: 0.1816, Validation Loss: 0.7980

```
Epoch [57/75], Training Loss: 0.1752, Validation Loss: 0.7960
Epoch [58/75], Training Loss: 0.1722, Validation Loss: 0.7926
Epoch [59/75], Training Loss: 0.1688, Validation Loss: 0.8316
Epoch [60/75], Training Loss: 0.1700, Validation Loss: 0.8164
Epoch [61/75], Training Loss: 0.1729, Validation Loss: 0.8212
Epoch [62/75], Training Loss: 0.1668, Validation Loss: 0.8493
Epoch [63/75], Training Loss: 0.1658, Validation Loss: 0.8635
Epoch [64/75], Training Loss: 0.1543, Validation Loss: 0.8433
Epoch [65/75], Training Loss: 0.1567, Validation Loss: 0.8254
Epoch [66/75], Training Loss: 0.1546, Validation Loss: 0.8731
Epoch [67/75], Training Loss: 0.1520, Validation Loss: 0.8600
Epoch [68/75], Training Loss: 0.1507, Validation Loss: 0.8495
Epoch [69/75], Training Loss: 0.1504, Validation Loss: 0.8276
Epoch [70/75], Training Loss: 0.1506, Validation Loss: 0.8634
Epoch [71/75], Training Loss: 0.1438, Validation Loss: 0.8609
Epoch [72/75], Training Loss: 0.1438, Validation Loss: 0.8186
Epoch [73/75], Training Loss: 0.1473, Validation Loss: 0.8702
Epoch [74/75], Training Loss: 0.1670, Validation Loss: 0.8136
Epoch [75/75], Training Loss: 0.1493, Validation Loss: 0.8392
```



```
Final Validation MSE: 0.8392
Final Validation RMSE: 0.9161
Model Complexity (Total Parameters): 12033
Total Training Time: 1.63 seconds
```

```
In [ ]: #2c
# Import necessary Libraries for the deep Learning model
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import time
```

```
# Load the housing dataset
data_source = "https://raw.githubusercontent.com/HamedTabkhi/Intro-to-ML/refs/heads/main/housing.csv"
raw_data = pd.read_csv(data_source)

# Convert categorical variables to numeric codes
for column in raw_data.select_dtypes(include=['object']).columns:
    raw_data[column] = raw_data[column].astype('category').cat.codes

# Initialize and apply one-hot encoding
cat_encoder = OneHotEncoder(sparse_output=False, handle_unknown='ignore')
cat_columns = raw_data.select_dtypes(include=['category', 'object']).columns
encoded_cats = cat_encoder.fit_transform(raw_data[cat_columns])
encoded_cats = pd.DataFrame(encoded_cats, columns=cat_encoder.get_feature_names_out)
processed_data = pd.concat([raw_data.drop(cat_columns, axis=1), encoded_cats], axis=1)

# Define and extract target variable
price_col = "SalePrice"
if price_col not in processed_data.columns:
    price_col = processed_data.columns[-1]

# Split features and target
house_features = processed_data.drop([price_col], axis=1)
house_prices = processed_data[price_col].values.reshape(-1, 1)

# Clean and scale the feature data
house_features = house_features.apply(pd.to_numeric, errors='coerce').fillna(0)
feature_scaler = StandardScaler()
scaled_features = feature_scaler.fit_transform(house_features)

# Split data into training and validation sets
train_x, val_x, train_y, val_y = train_test_split(scaled_features, house_prices, test_size=0.2, random_state=42)

# Convert to PyTorch tensors
train_x_tensor = torch.tensor(train_x, dtype=torch.float32)
train_y_tensor = torch.tensor(train_y, dtype=torch.float32)
val_x_tensor = torch.tensor(val_x, dtype=torch.float32)
val_y_tensor = torch.tensor(val_y, dtype=torch.float32)

# Create data Loaders for batch processing
train_data = TensorDataset(train_x_tensor, train_y_tensor)
val_data = TensorDataset(val_x_tensor, val_y_tensor)
train_loader = DataLoader(dataset=train_data, batch_size=32, shuffle=True)
val_loader = DataLoader(dataset=val_data, batch_size=32, shuffle=False)

# Define deep neural network for house price prediction
class DeepHousingNet(nn.Module):
    def __init__(self, input_size):
        super(DeepHousingNet, self).__init__()
        # Deeper architecture with gradually decreasing layer sizes
        self.input_layer = nn.Linear(input_size, 256)
        self.hidden1 = nn.Linear(256, 128)
        self.hidden2 = nn.Linear(128, 64)
        self.hidden3 = nn.Linear(64, 32)
        self.output_layer = nn.Linear(32, 1)
```

```
def forward(self, x):
    # Forward pass with ReLU activation
    x = torch.relu(self.input_layer(x))
    x = torch.relu(self.hidden1(x))
    x = torch.relu(self.hidden2(x))
    x = torch.relu(self.hidden3(x))
    x = self.output_layer(x)
    return x

# Initialize model and training components
input_dim = train_x.shape[1]
model = DeepHousingNet(input_dim)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model.to(device)
loss_fn = nn.MSELoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Training settings
num_epochs = 75
train_losses = []
val_losses = []
start_time = time.time()

# Training Loop
for epoch in range(num_epochs):
    # Training phase
    model.train()
    epoch_loss = 0.0
    for batch_x, batch_y in train_loader:
        # Move batch to device and compute predictions
        batch_x = batch_x.to(device)
        batch_y = batch_y.to(device)

        # Forward and backward passes
        optimizer.zero_grad()
        predictions = model(batch_x)
        loss = loss_fn(predictions, batch_y)
        loss.backward()
        optimizer.step()
        epoch_loss += loss.item() * batch_x.size(0)

    # Calculate average training loss
    avg_train_loss = epoch_loss / len(train_loader.dataset)
    train_losses.append(avg_train_loss)

    # Validation phase
    model.eval()
    val_loss = 0.0
    with torch.no_grad():
        for batch_x, batch_y in val_loader:
            batch_x = batch_x.to(device)
            batch_y = batch_y.to(device)
            predictions = model(batch_x)
            loss = loss_fn(predictions, batch_y)
            val_loss += loss.item() * batch_x.size(0)
```

```
# Calculate average validation loss
avg_val_loss = val_loss / len(val_loader.dataset)
val_losses.append(avg_val_loss)

print(f"Epoch [{epoch+1}/{num_epochs}], Training Loss: {avg_train_loss:.4f}, Va

# Calculate total training time
total_time = time.time() - start_time

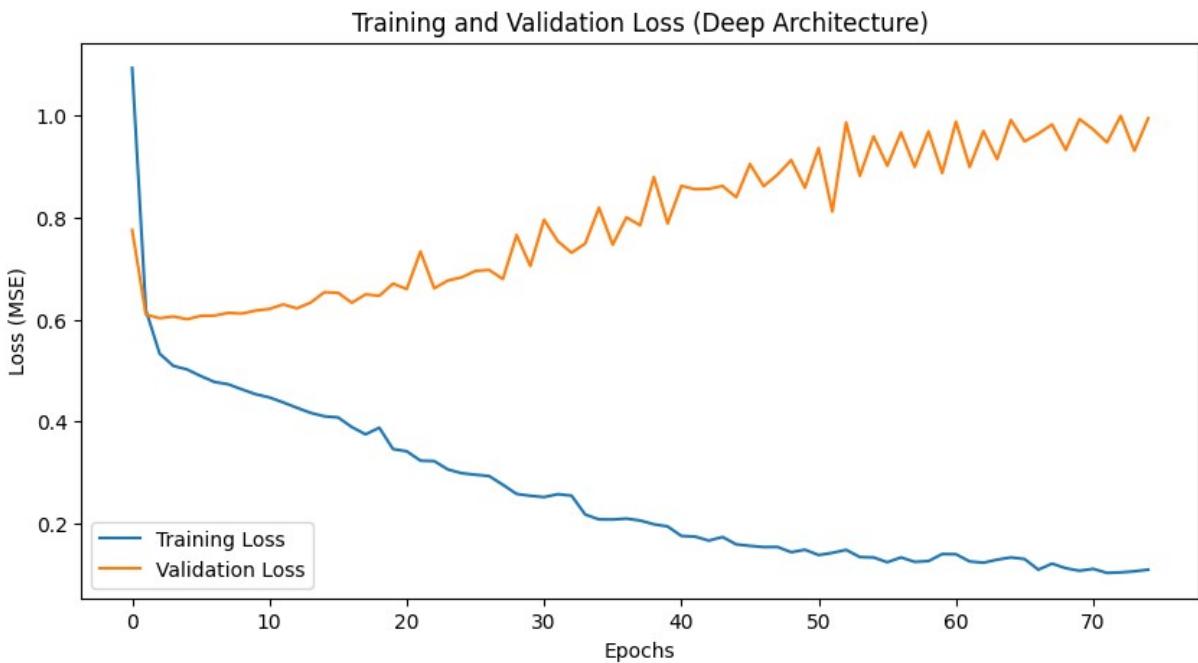
# Plot training and validation loss
plt.figure(figsize=(10, 5))
plt.plot(train_losses, label='Training Loss')
plt.plot(val_losses, label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss (MSE)')
plt.title('Training and Validation Loss (Deep Architecture)')
plt.legend()
plt.show()

# Final model evaluation
model.eval()
with torch.no_grad():
    val_x_tensor = val_x_tensor.to(device)
    val_y_tensor = val_y_tensor.to(device)
    final_preds = model(val_x_tensor)
    final_mse = loss_fn(final_preds, val_y_tensor).item()
    final_rmse = np.sqrt(final_mse)

# Print final metrics
print(f"Final Validation MSE: {final_mse:.4f}")
print(f"Final Validation RMSE: {final_rmse:.4f}")
print(f"Model Complexity (Total Parameters): {sum(p.numel() for p in model.parameters())}")
print(f"Total Training Time: {total_time:.2f} seconds")
```

Epoch [1/75], Training Loss: 1.0938, Validation Loss: 0.7758  
Epoch [2/75], Training Loss: 0.6183, Validation Loss: 0.6098  
Epoch [3/75], Training Loss: 0.5329, Validation Loss: 0.6025  
Epoch [4/75], Training Loss: 0.5089, Validation Loss: 0.6058  
Epoch [5/75], Training Loss: 0.5019, Validation Loss: 0.6006  
Epoch [6/75], Training Loss: 0.4889, Validation Loss: 0.6071  
Epoch [7/75], Training Loss: 0.4773, Validation Loss: 0.6076  
Epoch [8/75], Training Loss: 0.4728, Validation Loss: 0.6130  
Epoch [9/75], Training Loss: 0.4627, Validation Loss: 0.6117  
Epoch [10/75], Training Loss: 0.4532, Validation Loss: 0.6175  
Epoch [11/75], Training Loss: 0.4470, Validation Loss: 0.6207  
Epoch [12/75], Training Loss: 0.4372, Validation Loss: 0.6296  
Epoch [13/75], Training Loss: 0.4266, Validation Loss: 0.6218  
Epoch [14/75], Training Loss: 0.4167, Validation Loss: 0.6332  
Epoch [15/75], Training Loss: 0.4099, Validation Loss: 0.6536  
Epoch [16/75], Training Loss: 0.4077, Validation Loss: 0.6524  
Epoch [17/75], Training Loss: 0.3887, Validation Loss: 0.6328  
Epoch [18/75], Training Loss: 0.3745, Validation Loss: 0.6493  
Epoch [19/75], Training Loss: 0.3874, Validation Loss: 0.6465  
Epoch [20/75], Training Loss: 0.3456, Validation Loss: 0.6704  
Epoch [21/75], Training Loss: 0.3412, Validation Loss: 0.6597  
Epoch [22/75], Training Loss: 0.3229, Validation Loss: 0.7334  
Epoch [23/75], Training Loss: 0.3220, Validation Loss: 0.6611  
Epoch [24/75], Training Loss: 0.3055, Validation Loss: 0.6766  
Epoch [25/75], Training Loss: 0.2982, Validation Loss: 0.6828  
Epoch [26/75], Training Loss: 0.2952, Validation Loss: 0.6953  
Epoch [27/75], Training Loss: 0.2926, Validation Loss: 0.6973  
Epoch [28/75], Training Loss: 0.2755, Validation Loss: 0.6793  
Epoch [29/75], Training Loss: 0.2575, Validation Loss: 0.7658  
Epoch [30/75], Training Loss: 0.2538, Validation Loss: 0.7050  
Epoch [31/75], Training Loss: 0.2513, Validation Loss: 0.7958  
Epoch [32/75], Training Loss: 0.2570, Validation Loss: 0.7537  
Epoch [33/75], Training Loss: 0.2542, Validation Loss: 0.7311  
Epoch [34/75], Training Loss: 0.2170, Validation Loss: 0.7492  
Epoch [35/75], Training Loss: 0.2075, Validation Loss: 0.8195  
Epoch [36/75], Training Loss: 0.2073, Validation Loss: 0.7468  
Epoch [37/75], Training Loss: 0.2090, Validation Loss: 0.8005  
Epoch [38/75], Training Loss: 0.2053, Validation Loss: 0.7849  
Epoch [39/75], Training Loss: 0.1979, Validation Loss: 0.8796  
Epoch [40/75], Training Loss: 0.1937, Validation Loss: 0.7886  
Epoch [41/75], Training Loss: 0.1751, Validation Loss: 0.8622  
Epoch [42/75], Training Loss: 0.1737, Validation Loss: 0.8558  
Epoch [43/75], Training Loss: 0.1659, Validation Loss: 0.8564  
Epoch [44/75], Training Loss: 0.1729, Validation Loss: 0.8621  
Epoch [45/75], Training Loss: 0.1585, Validation Loss: 0.8400  
Epoch [46/75], Training Loss: 0.1556, Validation Loss: 0.9050  
Epoch [47/75], Training Loss: 0.1532, Validation Loss: 0.8617  
Epoch [48/75], Training Loss: 0.1534, Validation Loss: 0.8842  
Epoch [49/75], Training Loss: 0.1431, Validation Loss: 0.9130  
Epoch [50/75], Training Loss: 0.1479, Validation Loss: 0.8586  
Epoch [51/75], Training Loss: 0.1375, Validation Loss: 0.9368  
Epoch [52/75], Training Loss: 0.1418, Validation Loss: 0.8119  
Epoch [53/75], Training Loss: 0.1477, Validation Loss: 0.9870  
Epoch [54/75], Training Loss: 0.1336, Validation Loss: 0.8820  
Epoch [55/75], Training Loss: 0.1328, Validation Loss: 0.9596  
Epoch [56/75], Training Loss: 0.1235, Validation Loss: 0.9020

```
Epoch [57/75], Training Loss: 0.1328, Validation Loss: 0.9672
Epoch [58/75], Training Loss: 0.1241, Validation Loss: 0.8995
Epoch [59/75], Training Loss: 0.1257, Validation Loss: 0.9693
Epoch [60/75], Training Loss: 0.1395, Validation Loss: 0.8876
Epoch [61/75], Training Loss: 0.1391, Validation Loss: 0.9883
Epoch [62/75], Training Loss: 0.1250, Validation Loss: 0.8994
Epoch [63/75], Training Loss: 0.1225, Validation Loss: 0.9699
Epoch [64/75], Training Loss: 0.1284, Validation Loss: 0.9146
Epoch [65/75], Training Loss: 0.1329, Validation Loss: 0.9917
Epoch [66/75], Training Loss: 0.1297, Validation Loss: 0.9496
Epoch [67/75], Training Loss: 0.1088, Validation Loss: 0.9648
Epoch [68/75], Training Loss: 0.1206, Validation Loss: 0.9829
Epoch [69/75], Training Loss: 0.1116, Validation Loss: 0.9329
Epoch [70/75], Training Loss: 0.1067, Validation Loss: 0.9935
Epoch [71/75], Training Loss: 0.1104, Validation Loss: 0.9730
Epoch [72/75], Training Loss: 0.1025, Validation Loss: 0.9476
Epoch [73/75], Training Loss: 0.1035, Validation Loss: 0.9999
Epoch [74/75], Training Loss: 0.1055, Validation Loss: 0.9312
Epoch [75/75], Training Loss: 0.1089, Validation Loss: 0.9950
```



Final Validation MSE: 0.9950  
Final Validation RMSE: 0.9975  
Model Complexity (Total Parameters): 46593  
Total Training Time: 1.63 seconds